



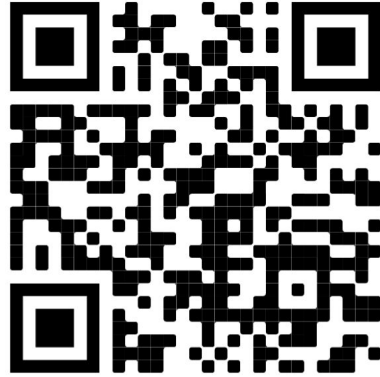
Introduction to Machine Learning

BOSTON UNIVERSITY
MACHINE INTELLIGENCE
COMMUNITY

Darcy Meyer
09/30/2021



Attendance




<https://forms.gle/1YDi83J3rvaWUanM8>

What is Machine Learning?

Machine Learning is the field of study that gives computers the capability to learn without being explicitly programmed




What is Deep Learning



Deep learning is a subset of machine learning

What is Deep Learning

Deep learning learns from data using a class of functions known as Neural Networks




A neural network maps an input to an output

Applications of Deep Learning

1. Cool things using deep learning


- a. Computer Vision
 - i. Tesla recognizing items on a street
- b. Text generation
 - i. An algorithm was trained to create a similar Shakespeare piece
- c. Image recognition
 - i. Classifying what a certain picture contains
 - ii. Facebook photo tagging
- d. Many more...






Learning from data

We have some data D



Make an assumption about D




$$y = b + mx$$

$$\hat{f} = \theta_0 + \theta_1 x$$

What is learning?

The approximation of some unknown function f based on some data D .



$$f : X \rightarrow Y$$

$$\hat{f} = \theta_0 + \theta_1 x$$

How do we set the parameters?


How do we know what assumptions to make?



Linear Regression

We have some data D

The approximation of some unknown function f based on D .




$$f : X \rightarrow Y \text{ or } y = b + mx$$

$$\hat{f} = \theta_0 + \theta_1 x$$

Initially, the theta parameters act as an arbitrary estimate for the parameters we are trying to learn.

Initial Guess

In this example we can instantiate our guesses (θ_0, θ_1) to be values close to zero. For example, the guesses in the example below are -0.01 and 0.01.



First Guess/Initialization:


$$\hat{f} = -0.01 + 0.01x$$

Computing the cost

In order to train our neural network, we need some way to tell us how far off its estimate was from the actual value.

We define the cost function, $J(\hat{y}, y)$ as the sum of losses, $\sum_{i=0}^m L(\hat{y}, y)$

- a. Loss = Error for a single training example
- b. Cost = Sum of all Losses
- c. y is our actual point
- d. \hat{y} is our estimated point




Compute the cost

Calculate difference between predicted output and actual data


$$J(\hat{y}, y) = \frac{1}{2n} \sum_i^n (y_i - \hat{y}_i)^2$$

Where i is the i th training example and n is the number of training examples

Intuition: if \hat{y} is far away from y , the square will be large



Cost function for gradient descent



Update Rule

To find the slope, we compute the derivative of the cost (gradient) with respect to a single parameter.

Also written $\nabla J(\theta)$

$$\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$$

Scalar learning rate

Individual weights

Cost/objective/loss function

Vector of weights


The diagram illustrates the weight update rule $\theta_j := \theta_j - \alpha \frac{\partial}{\partial \theta_j} J(\theta)$. Annotations include: 'Scalar learning rate' pointing to α ; 'Individual weights' pointing to the θ_j on the left; 'Cost/objective/loss function' pointing to $J(\theta)$; 'Vector of weights' pointing to θ in $J(\theta)$; and 'Also written $\nabla J(\theta)$ ' with a bracket pointing to the derivative term $\frac{\partial}{\partial \theta_j} J(\theta)$.



Linear Classifier

A Single Neuron


- Detects electrical pulses as inputs, and outputs one (or more) signal based on the inputs



Modeling a Single Neuron

- Want a function that takes n inputs and produces one output
 - Choose a simple linear function: $y = a_0 + a_1x_1 + a_2x_2 + \dots + a_nx_n$
- We want the output to be in a certain range (e.g. between 0 and 1)
 - Apply an "activation function" to y , e.g. $f(y) = 1/(1+e^{-y})$ (sigmoid function)

This is called a "perceptron"



What can we use it for? A Linear Classifier

Linear Classifier: given some input data, output "yes" or "no"

- Input data is a set of numbers x_1, x_2, \dots, x_n
- $z = a_0 + a_1x_1 + a_2x_2 + \dots + a_nx_n$ (where a_i are the weights of this model)
- Apply an activation function $f(z) = 1/(1+e^{-z})$ (squeeze outputs between 0 and 1)
- Interpret ~ 1 as "yes" and ~ 0 as "no"


Note that y is a dot product between \mathbf{a} and \mathbf{x} (first element of \mathbf{x} (x_0) is 1)




Activation Functions

Activation Functions model nonlinear data by taking inputs and comparing them to a threshold. This allows us to model non-linear data.


Sigmoid: output is
between 0,1



Tanh: output is
between -1,1




ReLu: output is
positive real numbers




Linear Classifier - Example

Parameters: $a_0 = -2$, $a_1 = 3$, $a_2 = 4$



- $z = a_0 + a_1x_1 + a_2x_2$
- $y = f(z) = 1/(1+e^{-z})$



- y around 1 \rightarrow red
- y around 0 \rightarrow blue

Training a Linear Classifier

Given some examples of labeled data points, how do we find a_0, a_1, \dots ?



Gradient Descent

- Start with random parameters, use them to predict an output
- Compare the predicted output to the correct output by computing a cost
- Find the partial derivative of the cost with respect to each parameter
- Update the parameters to decrease the cost
- Repeat many times



Limitations of Linear Classifiers

- What if we want to model something more complicated?
 - Data isn't linearly separable
 - Want to approximate any arbitrary function





Dogs?

An abstract graphic on the left side of the slide depicts a neural network. It features several circular nodes of varying sizes, colored in a gradient from light orange at the top to light purple at the bottom. These nodes are interconnected by smooth, flowing lines that follow a similar color gradient, creating a sense of dynamic movement and connectivity. The overall shape of the graphic is roughly circular, with lines looping and connecting the nodes in a complex, organic pattern.

Neural Networks

Neural Networks

Use outputs from some perceptrons as inputs to more perceptrons



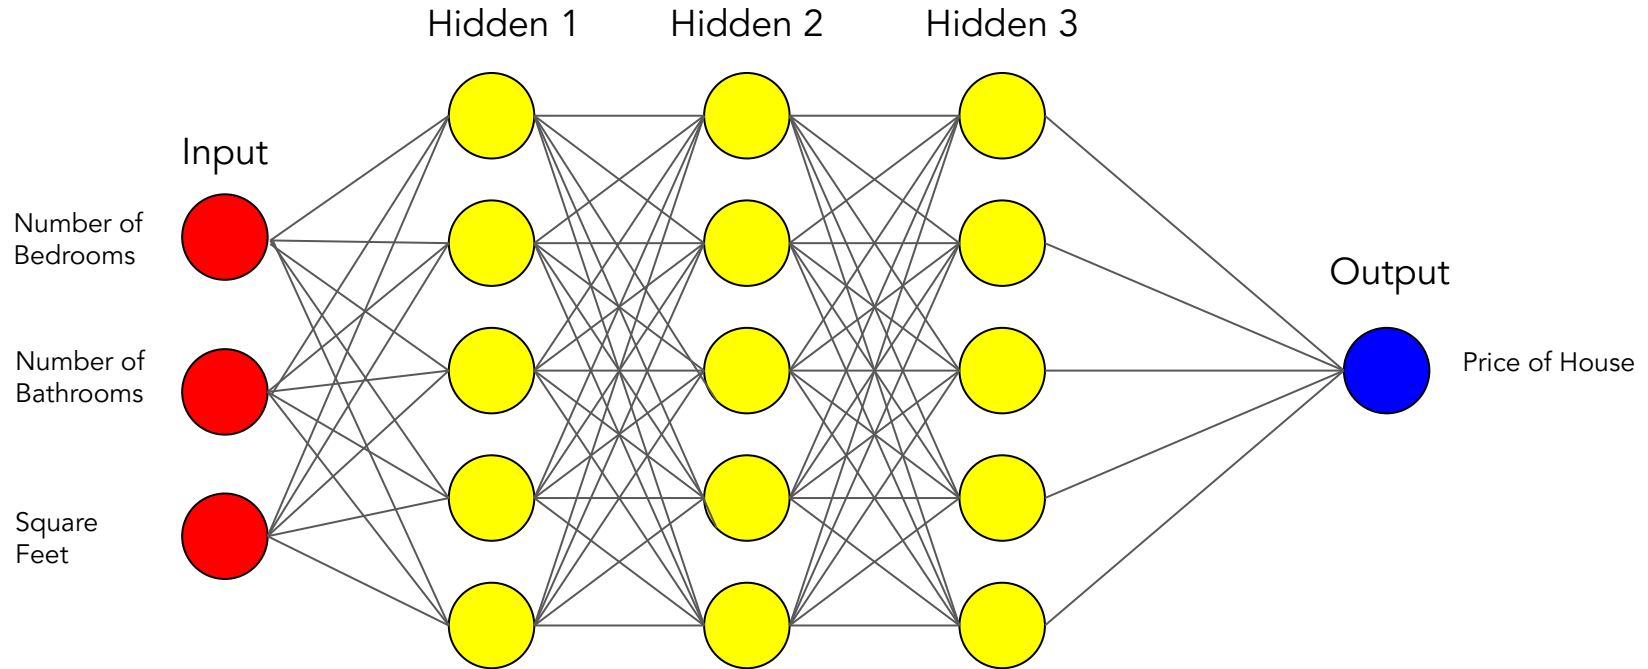
- Each single perceptron can detect one simple thing (one “feature”)
- Many perceptrons assembled together can detect complex things
- This is a “multilayer perceptron” or neural network



Steps to Train a NN

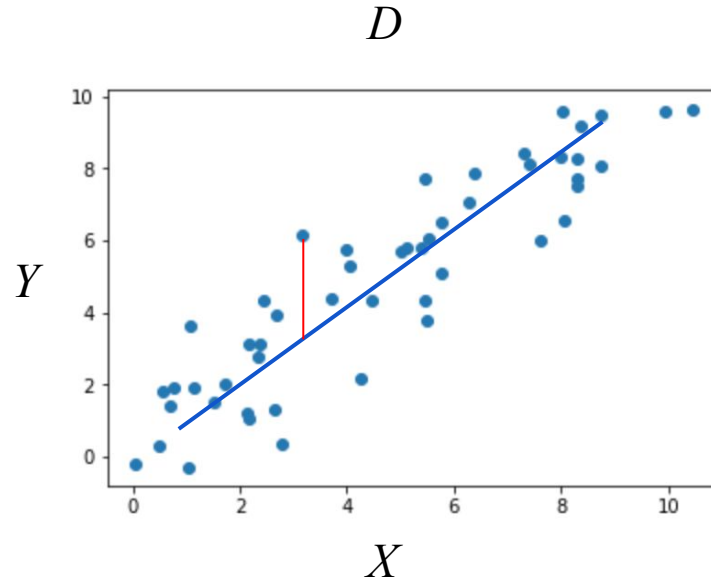

Forward propagation

Push example through the network to get a predicted output




Compute the cost

Calculate difference between predicted output and actual data



Compute the cost

Calculate difference between predicted output and actual data



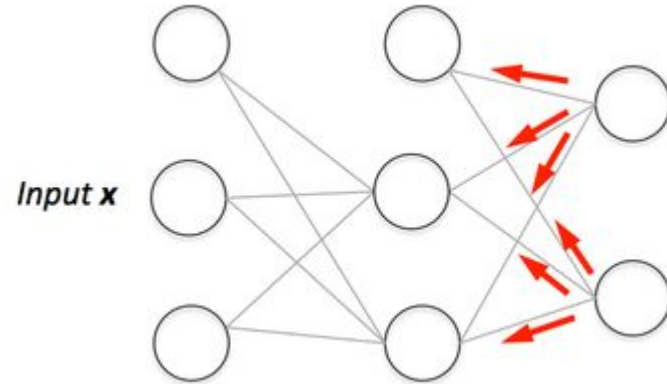
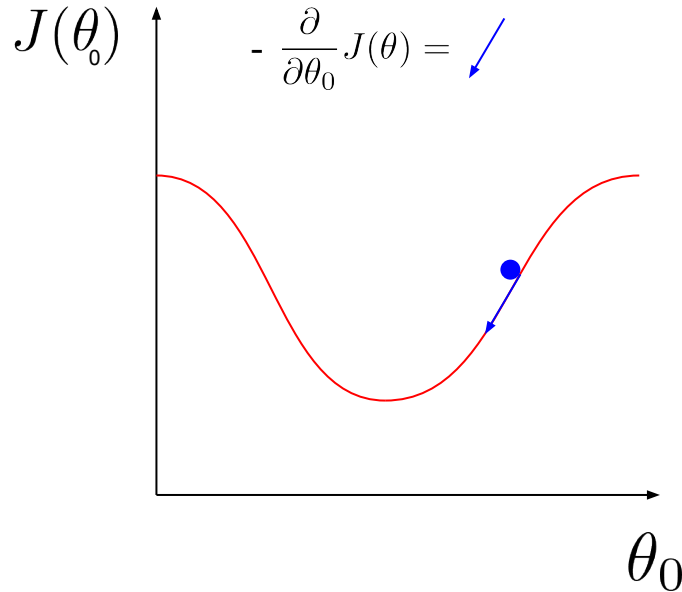
$$J(\theta) = \frac{1}{2m} \sum_i^m (y_i - \hat{y}_i)^2$$

Where i is the i th training example and m is the number of training examples



Backward propagation - "Update"

Push back the derivative of the error and apply to each weight, such that next time it will result in a lower error





Programming Exercise

<https://bit.ly/2Y6RMhS>



Eboard positions available!

<https://forms.gle/aV12v3iJVMnRb1xo6>